

Name: Huyen Nguyen

# Report for Distributed Shell

## Design

a) What programs/scripts you ran and what they did (use pseudo-code)

I used `measure_time` (compiled from `measure_time.c`) and `test_script.sh`.

### Pseudocode for `test_script.sh`

```
# Measure the amount of time required (the latency, in milliseconds) to setup a connection to the server, authenticate, and tear it down
```

```
for i = 0 to 10
```

```
    call measure_time and execute an empty command
```

```
# Measure the maximum throughput (in bits per second) from the server to the client
```

```
for i = 0 to 10
```

```
    create file$i.txt of size  $(4 * i)$  bytes
```

```
    scp the file to the server
```

```
    call measure_time and execute "cat file$i.txt" command
```

```
# Cleaning up
```

```
Shut down the server by making client send "exit" command to the server
```

### Pseudocode for `measure_time.c`

```
start = get current time
```

```
child_pid = Fork child process
```

```
If (child_pid == 0)
```

```
    Call rm using execvp()
```

```
    Print error // If execvp returns, it must have failed
```

```
Else
```

```
    Wait for child process
```

```
    Sync
```

```
    end = get current time
```

```
    print(start-end)
```

```
    Return child status
```

## b) How many runs you performed

I executed the empty command 10 times to measure the amount of time required (the latency, in milliseconds) to setup a connection to the server, authenticate, and tear it down. Also, in order to measure the maximum throughput (in bits per second) from the server to the client, I created 10 files of different sizes and use `measure_time` to measure the duration to transfer each of them.

## c) How you recorded your data

I used `measure_time.c` to record time. Specifically, for every file, I recorded the time right before running `rm` (1). Then, I created a child process to run `rm` and made the parent process wait for it. Then, when the child process is finished, I used `sync()`, and then measure the current time again (2). Finally, the time difference between (2)-(1) is the time taken to move the file to the dumpster.

## d) What the system conditions were like

The experiment was conducted on MacBook Pro (Retina, 13-inch, Early 2015) with macOS High Sierra. I made sure that the scripts run on the Linux VM for this course. The reason why I decided to run the tests on my actual OS because I experienced problems with `gettimeofday()` in VMs in OS class.

# Results

## Network Performance

	Time (seconds)	Time (microseconds)	
	1.008711	1008711	
	1.017965	1017965	
	1.049441	1049441	
	1.012894	1012894	
	1.029309	1029309	
	1.001871	1001871	
	1.003176	1003176	
	1.012418	1012418	
	1.022916	1022916	
	1.003716	1003716	
Mean	1.0162417	1016241.7	
Standard Deviation	0.01467093	14670.9305	

Figure 1: The amount of time required (the latency, in seconds) to setup a connection to the Amazon EC2 server, authenticate, and tear it down (10 iterations)

	File Size (in Bytes)	Time (seconds)	Time (microseconds)	
	4	1.004841	1004841	
	8	1.026714	1026714	
	12	1.004077	1004077	
	16	1.004343	1004343	
	20	1.000114	1000114	
	24	1.005963	1005963	
	28	0.998026	998026	
	32	1.035983	1035983	
	36	1.025691	1025691	
	40	1.05265	1052650	
Mean		1.0158402	1015840.2	
Standard Deviation		0.018347663	18347.6629	

Figure 2: The amount of time required to transfer a file of various sizes from the Amazon EC2 server to the client




Figure 3: The linear equation with time (in seconds) and file size (in bytes) using data from figure 2 and the mean value from figure 1 as the y-intercept




Figure 4: The linear equation with time (in seconds) and file size (in bytes) using data from figure 2 only




Figure 4: The quadratic equation with time (in seconds) and file size (in bytes) using data from figure 2 only

## CPU and I/O Performance

### Local Machine

#### CPU Performance

```
> sysbench --test=cpu --cpu-max-prime=20000 run
```

WARNING: the --test option is deprecated. You can pass a script name or path on the command line without any options.

sysbench 1.1.0-27a5b99 (using bundled LuaJIT 2.1.0-beta3)

Running the test with following options:

Number of threads: 1

Initializing random number generator from current time

Prime numbers limit: 20000

Initializing worker threads...

Threads started!

CPU speed:

events per second: 551.23

Throughput:

events/s (eps): 551.2305

time elapsed: **10.0013s**

total number of events: 5513

Latency (ms):

min:	1.41
avg:	1.81
max:	30.54
95th percentile:	2.76
sum:	9992.30

Threads fairness:

events (avg/stddev):	5513.0000/0.00
execution time (avg/stddev):	9.9923/0.00

## I/O Performance

```
> sysbench --test=fileio --file-total-size=1G prepare
```

```
> sysbench --test=fileio --file-total-size=1G --file-test-mode=rndrw --max-time=30  
--max-requests=0 run
```

Throughput:

read: IOPS=4461.73 69.71 MiB/s (73.10 MB/s) (1)  
write: IOPS=2974.49 46.48 MiB/s (48.73 MB/s) (2)  
fsync: IOPS=9514.30

Latency (ms):

min:	0.00
avg:	0.06
max:	47.21
95th percentile:	0.19
sum:	29526.98 (3)

```
> sysbench --test=fileio --file-total-size=1G cleanup
```

## Amazon Linux Instance

### CPU Performance

```
> sysbench --test=cpu --cpu-max-prime=20000 run
```

WARNING: the --test option is deprecated. You can pass a script name or path on the command line without any options.

sysbench 1.1.0-e5c8052 (using bundled LuaJIT 2.1.0-beta3)

Running the test with following options:

Number of threads: 1

Initializing random number generator from current time

Prime numbers limit: 20000

Initializing worker threads...

Threads started!

CPU speed:

events per second: 340.73

Throughput:

events/s (eps): 340.7268

time elapsed: **10.0021s**

total number of events: 3408

Latency (ms):

min: 2.66

avg: 2.93

max: 3.23

95th percentile: 3.07

sum: 9997.33

Threads fairness:

events (avg/stddev): 3408.0000/0.00

execution time (avg/stddev): 9.9973/0.00

## I/O Performance

> sysbench --test=fileio --file-total-size=6G prepare

> sysbench --test=fileio --file-total-size=6G --file-test-mode=rndrw --max-time=30

--max-requests=0 run

Throughput:

read: IOPS=1328.00 20.75 MiB/s (21.76 MB/s) (1)

write: IOPS=885.33 13.83 MiB/s (14.51 MB/s) (2)

fsync: IOPS=2831.49

Latency (ms):

min: 0.00

avg: 0.20

max: 16.70

95th percentile: 0.73

sum: **29759.11 (3)**

> sysbench --test=fileio --file-total-size=6G cleanup

# Analysis

## Network Performance

The linear equation on figure 3 has an R-squared value of 0.0586, whereas the one on figure 4 has an R-squared value of 0.3244. This is indicative that the relation between the amount of time to transfer a file and the file size might not be linear. In fact, the quadratic equation on figure 5 has an R-squared value of 0.6823, which is greater than the linear equations' R-squared values on both figures 3 and 4. In fact, in order to predict the amount of time to transfer a file, other factors could have impacts, such as the OS, the processor speed, etc.

## CPU and I/O Performance

### Solve equation

$$\text{Local\_CPU} + \text{Local\_File\_I/O} = n * \text{Network} + \text{Remote\_CPU} + \text{Remote\_File\_I/O}$$

- Local\_CPU = 10.0013 (number is highlighted in blue on page 4)
- Local\_File\_I/O = Reading time + Writing time =  $(1) \times (3)/1000 + (2) \times (3)/1000 \approx 3597(\text{MB})$  (the numbers are highlighted in red on page 5). Also, I used  $(3)/1000$  since I want the unit to be in seconds so it can be multiplied by (1), whose unit is (MB/s).
- Network = 1.0162417 (number is the mean on figure 1)
- Remote\_CPU = 10.0021 (number is highlighted in blue on page 6)
- Remote\_File\_I/O = Reading time + Writing time =  $(1) \times (3)/1000 + (2) \times (3)/1000 \approx 1079(\text{MB})$  (the numbers are highlighted in red on page 6)

So  $n \approx 2478$  (MB) or  $2.478 \times 10^9$  (bytes).

However, due to variations between computer configurations (e.g. CPU, disk, etc.), the equation might not be the most accurate to model the relations among the variables. However, the fact that this yields a positive for  $n$  makes sense. In fact, for a decent computer, local transfer should be faster than remote transfer. However, the number seems quite big. The fact that I used different file sizes for local (1 GB) and server (6 GB) might have caused this.

## Other Comments

- The `--init-rng=on` option in the command `$ sysbench --test=fileio --file-total-size=16G --file-test-mode=rndrw --init-rng=on --max-time=30 --max-requests=0 run` is not present in the sysbench Github documentation. Therefore, I removed it to perform the experiments.
- The command `$ sysbench --test=fileio --file-total-size=16G prepare` gave this error `FATAL: Failed to write file! errno = 28 (No space left`



on device). Therefore, I had to use 6GB for the Amazon EC2 instance, and only 1GB for my computer based on the amount of free storage I have (using `$ df -h`).

## Sample Session

```
$ ./server
```

```
SOCKET CREATED SUCCESSFULLY! PORT NUMBER IS 1024
```

```
./server activating
```

```
PORT:1024
```

```
directory:/Users/huyennguyen/Documents/Documents/distributed-systems-wpi/distributed-shell
```

```
Socket ready to go! Accepting connections....
```

```
*****Listening for requests...*****
```

```
$ ./client -s 127.0.0.1 -c "ls"
```

```
Client output:
```

```
Done. Creating socket...
```

```
Created. Trying connection to server...
```

```
Sending username...
```

```
Username sent successfully!
```

```
876817783
```

```
Client: Encrypted password 87FXVT5PVcoXQ
```

```
Server said credentials were ok
```

```
Command is ls
```

```
Command sent successfully!
```

```
Makefile
```

```
Project2_DistributedShell.pdf
```

```
Project2_DistributedShell_slides.pdf
```

```
README.md
```

```
abc.txt
```

```
client
```

```
client.c
```

```
client.h
```

```
client.o
```

```
distributed-systems-wpi.pem
```

```
server
```

```
server.c
```

```
server.h
```

```
server.o
```

```
shared.h
```

```
test
```

```
test_file.0
```

test\_file.1  
test\_file.2  
test\_file.3  
test\_file.4  
test\_file.5  
test\_file.6  
test\_file.7  
test\_file.8  
test\_file.9  
ubuntu@54.149.136.240  
ubuntu@ec2-54-149-136-240.us-west-2.compute.amazonaws.com  
Finished printing command output from server.

**Server output:**

Received connection  
Here is the username: huyen  
876817783  
Here is the encrypted password: 87FXVT5PVcoXQ  
Here is the encrypted password on the server: 87FXVT5PVcoXQ  
Correct credentials!  
Here is the command: ls  
child with pid 5869 exited with status of 0

\*\*\*\*\*Listening for requests...\*\*\*\*\*

**\$ ./client -s 127.0.0.1 -c "exit"**

**Client output:**

Done. Creating socket...  
Created. Trying connection to server...  
Sending username...  
Username sent successfully!  
878078308  
Client: Encrypted password 87FXVT5PVcoXQ  
Server said credentials were ok  
Command is exit  
Command sent successfully!  
Finished printing command output from server.

**Server output:**

Received connection  
Here is the username: huyen  
878078308  
Here is the encrypted password: 87FXVT5PVcoXQ  
Here is the encrypted password on the server: 87FXVT5PVcoXQ  
Correct credentials!  
Here is the command: exit

Server exiting...

[1] 5831 killed ./server